Leveraging smart meter data to recognize home appliances

Markus Weiss[+#], Adrian Helfenstein[+], Friedemann Mattern[+], and Thorsten Staake[φ]

[+] Institute for Pervasive Computing	[φ] Information Management	[#] Engineering Systems Division
ETH Zurich	ETH Zurich	Massachusetts Institute of Technology

Abstract — **The worldwide adoption of smart meters that measure and communicate residential electricity consumption gives rise to the development of new energy efficiency services. Several particularly promising applications involve the disaggregation of individual appliances within a particular household in terms of their energy demand. In this paper we present an infrastructure and a set of algorithms that make use of smart meters together with smartphones to realize new energy efficiency services (such as itemized electricity bills or targeted energy saving advice). The smartphones, together with a novel filtering approach, much simplify the training process for appliances signature recognition. We also report on the performance of our system that was tested with 8 simultaneous devices, achieving recognition rates of 87%.**

Keywords—smart metering; non-intrusive load monitoring; energy monitoring; energy break down; electricity consumption

I. APLLIANCE-SPECIFIC CONSUMPTION FEEDBACK

The requirement to conserve energy, the modernization of the electrical grid infrastructure, and the growing share of electricity from intermittent sources (e.g., wind and photovoltaics) initiated a paradigm shift in the energy domain [1]. As a consequence, smart electricity meters are currently rolled out in many countries. Besides simplifying the meter reading processes for energy utilities, smart meters are seen as enablers for new services, flexible tariffs, and demand response programs in the context of the smart grid.

Smart electricity meters record much more detailed consumption information than classical electricity meters. In current deployments, energy consumption is logged and typically made available to energy utilities and consumers by feedback tools such as web sites or in-home displays. However, the information provided is often limited to the mere visualization of the consumption data, or, at best, augmented with efficiency scales or comparisons with average households. While this may already contribute to energy savings, it fails to unlock the full benefit for consumers, as it does not direct the attention to those appliances or actions that bear high saving potentials.

In this work we propose a scheme that leverages metering data by automatically analyzing the recorded consumption information to provide better-tailored energy feedback at no extra cost. It provides users with an appliance-specific consumption break down. Such device-level information is essential to establish the link between consumption and device utilization, to enable sophisticated energy efficiency services (e.g., targeted automated recommendations), and to reduce residential electricity consumption by enabling users to derive conservation measures.

Most approaches so far have focused on providing this device-level consumption information by deploying sensors at appliances or power outlets. However, this is costly and the installation of a large number of sensors imposes a high usage barrier. Other solutions are based on a single sensor only, but require technical expertise for their setup, a-priori knowledge of appliance power signatures, and a complex calibration by the user [1].

Our approach tries to overcome these challenges. To facilitate the appliance-specific breakdown, we extended the capabilities of an earlier prototype that connects a smart meter with a mobile phone [2]. We disaggregate the recorded total load to device-level consumption information by applying data analytics to the electricity consumption data that is gathered by the smart meter and by making use of a measurement feature implemented as part of a mobile phone application. This not only enables more meaningful consumption feedback and increases users' energy literacy, but also leverages the added value of smart metering.

II. RELATED WORK

Based on the number of sensors used to gather device-level electricity consumption information, existing solutions for appliance load monitoring can be classified into two domains: multi-sensor systems and single-point sensor systems.

Multi-sensor approaches typically require a current sensor to be installed in-line with every device. To monitor the whole house, this device-level information is then aggregated at a central point. Commercially available solutions typically come in the form of smart power outlets. They measure the power consumption at the point where the load is caused and either visualize the data on a small display directly attached to the unit or propagate the consumption values wirelessly to a central display. A drawback of these systems is that they typically give a rather technical feedback and fail to integrate the consumption in a bigger picture that makes it more tangible for users. As mentioned above, such approaches typically require high monetary investment and high user effort to setup the feedback system.

Single sensor approaches are typically subsumed under the concept of Nonintrusive Appliance Load Monitoring (NALM). The initial work dates back to the 1980s, where Hart [3] tried first to match a-priori known appliance signatures in the overall power signal by using real and reactive power measurements at a rate of 1Hz. The concept proved to be effective in several field tests – at that time especially for larger loads – and paved ground for various other work based on this principle. Norford and Leeb, for example, introduced transient event detection at high sampling rates to disaggregate devices with similar power consumption [4], and follow-up work by Laughman et al. [5] explained how to use current harmonics to further disaggregate continuous variable loads. A variant of Hart's scheme deals with the separation of simultaneous on/off events of appliances [6].

Other work utilizes methods from artificial intelligence to disaggregate overall residential energy consumption data. Early approaches were typically bound to low-resolution data. Powers' [7] rule-based algorithm tries to analyze the energy consumption at a low sampling interval of 15 minutes. However, his approach is based on a large a-priori known reference database that requires monitoring of each appliance in the home for several days. Prudenzi disaggregated consumption data for large loads at the same sampling rate by using a neural network approach [8]. Ruzzelli et al. used a special purpose sensor that has to be installed at the circuit breaker. The consumption information is post-processed in an artificial neuronal network that requires a lengthy training process to disaggregate device level consumption [9]. Other rule-based work focuses on the possibility of differentiating between appliances with similar power consumption by taking into account their frequency of use [10] and using pattern recognition methods to disaggregate the overall electricity consumption into major energy end-uses [11].

More recent approaches deal with the analysis of data sampled at higher frequency. Statistical signature analysis has been used to infer the devices operating from the current and voltage waveforms [12]. Srinivasan combined harmonic signature analysis with neural networks and developed and tested several different classification models for signature extraction and device identification [13]. In contrast to these high frequency approaches that usually rely on special purpose sensors, Kolter et al. [14] recently investigated the possibility of load disaggregation using discriminative sparse coding based on hourly data.

Yet another idea has been explored in [15] and [16]. The authors combined two complementary approaches in a system that relies on a single sensor that can be plugged-in anywhere to the electric circuit. It then listens to detect unique noise changes and electromagnetic interference that occur through the switching of devices and through switch mode power supplies. The system can be used to infer about device operation which in combination with the data of an electricity meter can reveal the consumption of particular devices.

Summarizing the related work, existing systems can be characterized as follows: Multi-sensor approaches can rather easily achieve a consumption breakdown, but deploying a large number of sensors in the residential environment quickly leads not only to high cost but also to a discouraging high usage barrier [9]. In contrast, single sensor systems are easier to deploy but often rely on expensive custom hardware (e.g., for high sampling rates) and require either a priori knowledge about the household devices and their electrical characteristics, or they require a complex training phase involving the user where the system learns about the specific device characteristics. However, a-priori knowledge is difficult to obtain in a world of fast changing small appliances, and training procedures at the initial deployment are discouraging users and hinder adoption [10]. In addition, these approaches cannot take into account new devices that are introduced into the residential environment. Overall, we conclude that existing approaches fail to meet usability requirements that are essential for fast adoption.

III. SYSTEM OVERVIEW

Our system represents an integrated solution to identify the electricity consumption of household appliances from the data gathered by smart meters, which will be installed in large numbers in many countries over the next years. The approach builds on an earlier principle advocated by Hart [3]. It uses a single sensor and addresses remaining technical challenges (e.g., the recognition of smaller loads and overlapping on/off events of appliances) as well as some of the above-mentioned shortcomings with respect to usability. For this, we designed and developed a system that does not rely on custom hardware or complex training. In particular, we make use of a smartphone application, which much simplifies the appliance signature acquisition process because this is done as a side effect, invisible to users.

In the following, we first give an overview on the system architecture and its components. We then explain how residential appliances can be classified according to their characteristic electricity consumption.

A. Data Acquisition Architecture

One of the three main components of our system is an electricity meter that can measure the total electrical load of all attached devices in a household (Figure 1). The meter (we used model E750 by Landis + Gyr) logs the total consumption at a frequency of 1 sample per second. It has an integrated communication interface that is connected to a gateway, which is responsible for continuous data acquisition and storage from the electricity meter, and also for the handling of the incoming requests of the user interface. For that, the gateway consists of a web server (lighttpd, php), an SML parser, and a database (SQLite3). It is implemented on an embedded device based on a 600MHz CPU, 256MB storage of RAM and flash memory, and an Ethernet and WiFi-module for communication purposes. The third component of the system is the user interface that is implemented as a smartphone application. It provides users with real-time feedback on their electricity consumption. A detailed description of the design and the capabilities of the system can be found in [1, 2]. The communication between the three decoupled components is realized over http following the "Web of Things" paradigm [17]. This integrates physical resources, such as the meter and its measurements, seamlessly into the web. They can then be identified by URLs and accessed by the four basic HTTP commands through the RESTful-API provided by the gateway [2].

B. Classification of Residential Appliances

In the following, we explain how domestic appliances

Figure 1. Loosely coupled data acquisition architecture.

can be classified according to their characteristic load signatures based on the physical quantities (i.e., apparent power, reactive power, real power, and distortion power) measured by the smart meter. Depending on its characteristic electrical and electronic components, an appliance can be of resistive, inductive, or capacitive nature. For example, a standard light bulb is purely resistive whereas a vacuum cleaner is predominantly inductive. In general, incandescent appliances (e.g., kettle, light bulb) are mostly resistive (ohmic), motors (e.g., fans, heaters) predominantly inductive, and devices containing a power supply or electronic frequency converters (e.g., laptops) mainly capacitive.

Figure 2 illustrates exemplary power signatures at a sampling frequency of 1Hz for different appliance categories over different operation lengths. If the load is purely resistive, then the voltage and current are in phase (e.g., the iron (Figure 2 (left)). The reactive component Q of the apparent power is null, meaning all power is transferred to the load. A consumer with reactive components is either of type ohmic-inductive with a typical phase shift of $0 < \varphi < \pi$ between current and voltage or ohmic-capacitive characterized by a negative phase shift $0 > \varphi > -\pi$ (Figure 2 (middle and right)). In addition, in electrical networks there may exist non-sinusoidal currents and voltages (e.g., caused by inverters in switching events) that result in harmonics. These harmonics cause an additional reactive component, the so-called distortion power (Figure 3). In mathematical terms this can be expressed as:

$$|S| = \sqrt{P^2 + Q_{trans}^2 + D^2},$$

where S is the apparent power, P is the real power, Q the translative component, and D the distortive component of the total reactive power.

Based on its internal composition and its possible modes of operation (e.g., static, multi-level, or variable) an appliance imposes a characteristic load profile on the electric circuit. This signature depends on the relation of the different power components and can be used to discriminate between appliances when disaggregating the total consumption. Our prototype system measures these parameters either directly or indirectly. In addition to these physical quantities, the signature length, peak voltage, and current are also important in terms of the appliance signature.

IV. THE APPLISENSE ALGORITHM

The AppliSense algorithm uses consumption data gathered by the smart electricity meter to automatically break down the total consumption to device-level. In the following, we first outline the basic idea and concept of our system that pays particular respect to usability. We then explain how the

signature database on which the algorithm crucially depends is acquired and discuss some algorithm details.

A. Basic Concept

The electricity consumption of a household fluctuates over time based on the operation of individual devices used by the residents (see Figure 4). For example, switching on a light induces the depicted change in the load curve. Having a more detailed look on the consumption data, the figure shows that there exist intervals where the load remains more or less constant on a stable level. A black bar marks two of these levels. The difference in real power (dP) between these levels indicates the change in electricity consumption due to the operation of the light. Our system not only measures the total load of the household, but the load characteristics (i.e., apparent power, real power, etc.) of each of the three phases separately. This phase-level data allows us to split up the overall electricity to get an even more detailed view.

These considerations lead to the following key concept of AppliSense, which can recognize device-switching events in the load curve based on an appliance signature database.

First, identify time points where significant changes between two levels of power consumption in the load curve occur. Second, once such an edge is detected, compute the differences of the different physical quantities between these consecutive levels and classify the change as a potential appliance-switching event. And third, compare each of these differences with a known set of differences from an appliance signature database and map the edge to an individual device according to its load characteristics.

Figure 5 illustrates these steps. It shows the electricity consumption (red) at a certain time interval in which five load levels (black bars) were identified. For simplicity, only the real power is visualized in this example. From this we


Figure 3. Relation between the different power quantities that can be derived with our prototype.


Figure 2. Power signatures of three different residential devices from different appliance categories and for different operation periods.

can compute four deltas: dP1, dP2, dP3, and dP4. Each of these deltas corresponds to a potential on/off event of a device. The algorithm tries to match these with a known device signature from the database. For that, each entry dP_i in a column of the matrix on the left symbolizes a delta which was extracted from the load curve at time i. The operator represents a detector logic that compares the rows of the matrix to the signature vector with the known deltas. The resulting vector holds the best matching entry, in case a matching appliance could be identified. In the example, this means that at time instant two and three matching signatures of a known device (a turning on and a turning off event) are detected. However, no signature is matching the events at time instants one and four.

B. *User-friendly Singature Recording*

In contrast to other load disaggregation systems, which often discourage users by requiring a long training period or complex calibration, we wanted to develop a system that is easy to use. This is particularly important for the generation of the signature database that is used to identify an appliance power signature. With our approach it is not necessary to take signatures of every appliance in advance, but the signature database is established with simple means over time. For that, we equipped the user interface of the smartphone with a measurement functionality that allows users to identify the consumption of an individual appliance in a simple, explorative way while at the same time logging the signature in the

Figure 4. Key idea of the AppliSense algorithm.

Detected Edges   **Known Signatures Recognition**

Figure 5. Simplified overview of the required steps to recognize an appliance in the overall electrical load.

background, invisible to the user. This also allows easily integrating new appliances that are introduced at home. Whereas other systems need to completely recalibrate, our approach is able to incrementally acquire signatures and thus integrate new devices, which is crucial in a fast changing home environment.

The measurement process is illustrated in Figure 6. To measure the consumption of a device, users initialize the measurement by pressing the start button on the user interface. They are then asked to turn the device being measured either on or off. Within a few seconds, the system then computes the result based on the measurement algorithm [18]. If desired, users can further personalize the measurement (e.g., picture, name, category, etc.) and store the device characteristics in the inventory of the mobile phone application.

During the measurement, the signature acquisition process runs in the background (see Figure 7; only real power depicted for clarity reasons). It logs the whole appliance signature (e.g., change in apparent, reactive, and distortion power, etc.). In addition, the algorithm automatically classifies whether an on (dP>0) or off (dP<0) switching event has occurred. AppliSense uses this information later as input knowledge. The idea of this approach is to systematically increase the number of signatures in the database while the system is being used. This leads to higher precision in recognizable operation events over time.

C. *Algorithm Design*

The AppliSense load disaggregation algorithm consists of six steps that are subsequently discussed in this section (Figure 8). It follows the early principles discovered by Hart, but much simplifies the signature acquisition process for users.

(1) Normalization and (2) Edge Detection: In power circuits, load-dependent voltage drops can occur (e.g., in reaction to a switching event of an appliance). From I = U/R and S = U×I for apparent power S and effective values of voltage U and current I, a quadratic relation arises: $S = U^2/R$.

Hence, voltage drops can lead to large differences in power consumption, which we have to account for by normalizing the power values to a constant voltage (of 230 V):
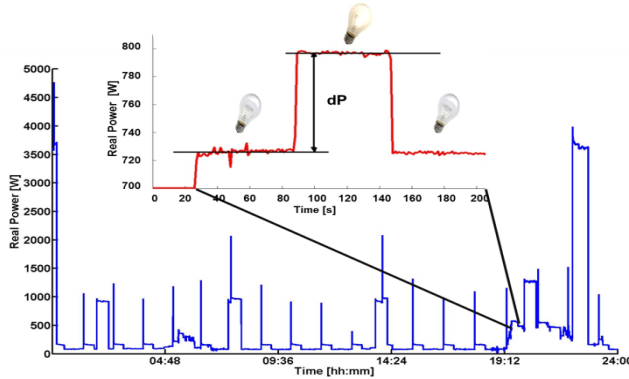

Figure 7. User-friendly signature acquisition process with the help of the measurement feature of the user interface.

Figure 6. User process to measure the power consumption of an individual appliance (e.g., an office flood light).

193

$$S'_n = \left(\frac{230}{U}\right)^2 \times S.$$

In order to identify edges in the recorded electricity consumption data that correspond to switching events of appliances, we use the normalized apparent power S'_n as input vector. The algorithm computes the absolute values of the differences between two consecutive values of normalized apparent power S'_n in the data series. If the absolute value of such a difference is larger than a predefined threshold f_th, then the value potentially belongs to an edge. However, there can be much more potential edges than appliance-switching events. The threshold f_th has to be robust to small changes due to noise on the electric line.

The leftmost plot in Figure 9 depicts the apparent power of a Nintendo Wii usage cycle over a time span of 180 seconds. The two distinctive edges are related to switching the game console on (at time step 22) and off (at time step 106). The figure also shows the relatively strong fluctuations in apparent power during the start phase of the device compared to the ones in standby (from 0s to 22s). The middle histogram of the figure depicts the difference of two subsequent apparent power values over the same time frame. We find larger changes in apparent power when turning the application on/off compared to times of operation or standby. We experimented with different thresholds and generally achieved best results applying a filter with a threshold f_th of 2VA. It removes a large number of intervals that do not correspond to a switching event (Figure 9 right). However, due to the transient behavior of the particular appliance, there persist some peaks (e.g., between 24 to 45 seconds) in the graph although no switching event occurred. In general, such oscillations during operation can be even stronger and more frequent which would result in a high number of spurious events. Applying a smoothing filter can help remove these false detections. However, it also bears the risk of cancelling out edges (typically small ones) that correspond to real switching events. Consequently, these switching events would not be identified and the operation would be missed.

In order to decrease the number of spurious events, we investigated different smoothing filters. We tested a median filter, a mean filter, a kernel-weighted average filter (Nadaraya-Watson filter with Gaussian kernel), and different combinations of these on the apparent power signal. An advantage of a median filter is the ability to remove outliers. However, periodic curves (e.g., sine, triangle, saw tooth, square, etc.) could be resistive. On the other hand, a mean filter, which computes equally weighted averages of a sliding

Figure 8. Overview of the six steps of the AppliSense algorithm.

window of values, has the ability to smoothen periodic oscillations, but may not always remove large outliers. Even worse, it might erase edges which correspond to an actual on/off switching of a device. A kernel-weighted average filter adds more complexity. We experimented with several kernel functions and observed best results when applying a Gaussian kernel [19]. It allows preserving edges while attenuating oscillations of the original signal. The extent to which the filter smoothens the signal is determined by the kernel bandwidth, which relates to the window size.

In order to evaluate the influence of the filters on the edge detection and to find the most appropriate combination of filtering, we simulated a typical household usage scenario over 30 minutes in a controlled lab environment. During that period appliances of different characteristics were used and 12 appliance switching events occurred. Table I shows the results when applying the different filters to the signal. The number in brackets corresponds to the window size/kernel bandwidth of the respective smoothing filter. The table displays the number of changes of apparent power values larger than 2VA, the achieved reduction gain compared to the original signal, and the number of missed true appliance on/off events. Overall, the original signal contained 709 changes in apparent power with a delta larger than 2VA.

Using a median filter or a mean filter alone reduces the number of potential edges by 74% and 70% respectively without missing a true device-switching event. A combination of mean and median filter achieves slightly better results (76%) at no extra cost in terms of computation complexity. The reason for this relatively small improvement is due to the fact that the possibility to remove outliers is constrained by the small window size. This filter parameter determines the extent to which the original signal is smoothened. However, increasing the value decreases the lower limit of loads that can be detected by AppliSense. The parameter of 5 was chosen as a trade off that enables filtering without precluding the recognition of smaller loads.

The performance of kernel smoothening strongly depends on the bandwidth of the kernel. The potential reduction varies between 35% and 94% depending on the kernel bandwidth. Adding a kernel filter to the smoothening median/mean strategy leads to higher reduction in potential edges (3% − 17%). From a bandwidth parameter of 60 on, we observe that the smoothening starts canceling out true switching events. Independent of the bandwidth, however, this approach increases the computational complexity significantly.

Overall, we achieved best results using a kernel filter. However, this comes at high computational cost due to the quadratic complexity of the filter. Hence, we decided to go for a more efficient solution that performs close to optimum. It combines a median filter that removes outliers with a mean filter that further smoothens the signal (see line 5 of Table I). The result of this smoothening strategy is illustrated in Figure 10. In our evaluation scenario, we used a notebook, several different lights, and a kettle to obtain the original power signal (red). The blue markers correspond to the 709 points in time where the absolute difference of two subsequent apparent power values is greater than 2VA. Applying a median filter of five followed by a mean filter of the same size results in the green markers. The reduction gain (75.5%) of the
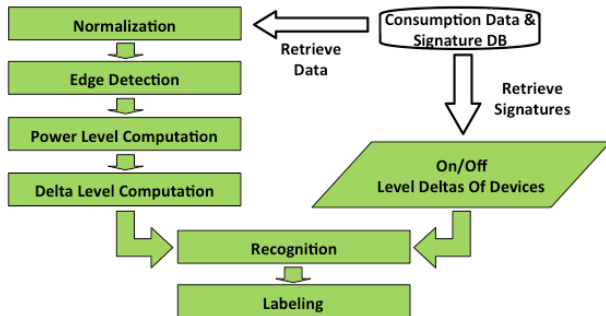
Figure 9. Apparent power for a duty cycle of a Nintendo Wii (left). Resulting absolute differences in apparent power (middle) and with a filter of 2VA (right).

filter can be seen by comparing the blue with the green markers. The edge detection interprets the remaining 174 green markers as a binary vector which indicates at position i that the smoothed estimate of the apparent power at time step i differs by more than 2VA from the value at position i-1. Hence, the measurement at time i belongs to a potential device-switching event.

TABLE I. COMPARISON OF DIFFERENT SMOOTHENING FILTERS

Filtering Method	ΔS > 2VA	Reduction	Missed
Median(5)	185	73.9%	0
Mean(5)	218	69.2%	0
Kernel(3)	459	35.3%	0
Kernel(100)	46	93.5%	0
Median(5), Mean(5)	174	75.5%	0
Median(5), Mean(5), Kernel(3)	151	78,7%	0
Median(5), Mean(5), Kernel(60)	78	89%	1
Median(5), Mean(5), Kernel(70)	52	92.7%	4

(3) Power Level and (4) Delta Level Computation: Having identified the relevant edges, the next step extracts power levels that connect two edges in the smoothened signal. From two consecutive power levels separated by an edge, the algorithm then extracts the delta vectors that are used for matching the edge to a particular device.

Each power level consists of a start and an end time, a vector with component-wise means of real, reactive, and distortion power for the first five measurements at the start and the last five measurements at the end of the interval (start mean (sm) vector and end mean (em) vector), and a three-by-five matrix which holds the original real, reactive, and distortion power values. The component-wise standard deviation of all power values is also calculated.


Figure 10. Application of mean/median filter to the original power signal.

From two consecutive power levels, the algorithm computes the difference vector for real, reactive, and distortion power. To take oscillations during start up and shut down of an appliance (e.g., due to heating up at the start of a kettle) into account, we not only calculate one difference vector for level i to i+1 (e.g., end of level i (em_i) – start of level i+1 (sm_{i+1})), but four difference vectors $d_{i,j}$ that include the start and the end values of both levels:

$$\overrightarrow{d_{s\_i,s\_i+1}} = \overrightarrow{sm_i} - \overrightarrow{sm_{i+1}},$$
$$\overrightarrow{d_{s\_i,e\_i+1}} = \overrightarrow{sm_i} - \overrightarrow{em_{i+1}},$$
$$\overrightarrow{d_{e\_i,s\_i+1}} = \overrightarrow{em_i} - \overrightarrow{sm_{i+1}}, \text{ and}$$
$$\overrightarrow{d_{e\_i,e\_i+1}} = \overrightarrow{em_i} - \overrightarrow{em_{i+1}}.$$

For each edge, we add these four vectors to a result matrix used as input for matching the device signatures in the next step.

(5) Recognition and (6) Labeling: The recognition part of the algorithm tries to match known appliance signatures $\overrightarrow{k_j}$ from the signature database with extracted delta vectors $\overrightarrow{d_i}$ obtained as a result in the previous step. In order to identify an appliance on/off event, we perform a nearest neighbor search in the two-dimensional dQ/dP space (see Figure 11).

First, the algorithm computes for every $\overrightarrow{d_i}$ its Euclidean distance to every $\overrightarrow{k_j}$ in the two-dimensional vector space. If this is smaller than a predefined value (r) of the length of $\overrightarrow{k_j}$ plus an oscillation value (osc), a potential matching is identified:

$$\|\overrightarrow{d_i} - \overrightarrow{k_j}\| < r \cdot \|\overrightarrow{k_j}\| + osc \begin{cases} if \ true, \overrightarrow{k_j} \ is \ a \ potential \ match \ for \ \overrightarrow{d_i} \\ if \ false, \overrightarrow{k_j} \ is \ not \ a \ match \ for \ \overrightarrow{d_i} \end{cases}$$

The oscillation term (osc) is the length of a vector which consists of the maximum of the standard deviation in the real power at level i or i+1 as first component, and of the maximum of the standard deviation in reactive power at level i or i + 1 as second component:

$$osc = \begin{pmatrix} \max(std(P \ at \ level \ i), \ std(P(at \ level \ i+1)) \\ \max(std(Q \ at \ level \ i), \ std(Q(at \ level \ i+1)) \end{pmatrix}$$

After this, every $\overrightarrow{d_i}$ contains a set of associated possible recognition candidates $\overrightarrow{k_j}$ from the signature database. Note that this set of possible associated recognitions could also be empty. In such a case, the corresponding $\overrightarrow{d_i}$ could not be related to a known signature. This could be caused for example by a detected edge which does not correspond to an appliance switching event, or by the non-existence of a corre-

sponding signature in the database that matches $\vec{d_i}$. Second, for each $\vec{k_j}$, a nearest neighbor match is performed over all potentially matching candidates $\vec{d_i}$ that have been associated with $\vec{k_j}$. Finally, the algorithm labels the load profile with the corresponding device names.

V. ALGORITHM EVALUATION AND LIMITATIONS

In order to analyze the performance of the AppliSense algorithm, we installed the whole system in a laboratory environment. For the evaluation, we used a controlled set of appliances which typically occur in a student's household. Table II provides an overview of the appliances, their real power consumption stated on the manufacturer label, their verified real power range in operation (measured by a separate power monitor), the appliance category (O for ohmic, I for ohmic-inductive, and C for ohmic-capacitive), and the real power that is obtained as part of the power signature using our smartphone application. All devices were connected to the same phase over the whole evaluation. Some of the appliances have power consumptions within the same range. However, if belonging to different categories, we should still be able to differentiate the corresponding events.

TABLE II. APPLIANCES USED FOR ALGORITHM EVALUATION

Appliance	Labeled Power	Power Range	Category	Consumption
Light bulb	75W	70W	O	70W
Kettle	2200W	1855 – 1933W	O	1900W
Heater	2000W	1619 – 1667W	O	1635W
CD player	13W	9 – 13W	I	3W
Fan	50W	45W	I	45W
Notebook	72W	30 – 35W	C	35W
Fluorescent lamp	35W	21 – 28W	C	25W
Wii	52W	10 – 45W	C	15W

During times when only a single appliance was active, the algorithm identified the on/off events of all devices except the CD player correctly. Every device was turned on and off at least three times. The edges caused by the CD player were not recognized neither when being turned on nor when being turned off. This can be explained through the limitations introduced by the filtering. Using a window size of 5 samples in our test scenario leads to a lower boundary of 10VA for edges that can be recognized. The CD player has a relatively high standby consumption of 6W compared to its 3 – 7W in operation. While the median filter does not influence the signal, the constant 3W during operation result in a


Figure 11. Device recognition: Comparing a detected power edge to a known signature in the Euclidian dQ/dP space.

step-wise increase of 0.6VA after the application of the mean filter. This increase is too small (<<2VA) to be detected as an event using the chosen median/mean filter.

Next, we combined the use of multiple devices in a random order. Although the CD player cannot be recognized, we operated it and other devices with unknown signatures from time to time to vary the base line consumption and to have more appliances concurrently running. Over a time span of several hours, we documented 80 switching events of which 77 were identified correctly.

Figure 12 shows a sample labeling output of the algorithm (for a simulated office environment). After the notebook has been turned on, different devices were concurrently used and a kettle was operated. However, the red circle highlights a moment at which the office lamp is turned on but the event is not detected. This is due to the oscillations caused by a device that was operating at the same time. A second problem (not depicted) occurred when operating the notebook. Due to the different battery levels, the power consumption had varied compared to the one registered in the appliance signature database. This led to the correct identification of the edge, but no appliance signature could be matched to the detected event.

These two examples outline limitations of the current implementation. Oscillations caused by operating devices can mask the consumption, especially of low power drawing appliances. This could in particular be a problem in larger households (e.g., family houses) with lots of appliances and activity. In addition, the algorithm cannot detect devices that do not have well-defined operation states but have a continuously changing consumption. This is due to the initial assumptions regarding the algorithm design and the tradeoff for relying on a single sensor system with a 1Hz sampling frequency. In the conducted laboratory study we observed that the appliance signatures recorded with the smartphone application were very reliable. That is, the delta vectors obtained with the measurement function when turning an appliance on/off are stable and reproducible over time. However, this may be different in a more dynamic home environment – there the algorithm may need several (slightly different) signatures per device to reliably recognize appliances.

Overall, the evaluation shows promising results. We generated 144 device-switching events in our test scenario. 16 of these came from devices with a consumption so small that


Figure 12. Labeled load curve as output of the AppliSense algorithm.

the filter canceled out the corresponding edges. When subtracting these events, the algorithm identified 125 out of the remaining 128 events correctly, which results in an overall recognition rate of about 90%. In practice this enables interesting applications, such as automated recommendations for a more economic use of electricity in households.

VI. CONCLUSIONS AND FUTURE WORK

We gave a detailed description and evaluation of a system that facilitates automatic recognition of switching events of electric appliances. In contrast to other existing approaches, our objective was to develop a system that achieves this by being unobtrusively integrated in users' life and without requiring a complex system setup or training. We achieved this by interconnecting components that are becoming ubiquitous in home environments: a smart meter and a smartphone. The signature database is established over time and also allows introducing new devices, which is important in a fast changing home environment. In particular, we achieve this as a side effect of a smartphone application, which much simplifies the appliance signature acquisition for users.

Applying data analytics to the gathered metering data allows the system to raise energy awareness by providing better-tailored energy feedback without the need for special purposed hardware. In combination with actuation capabilities, we can foresee this information to be used to automatically optimize energy consumption and hence increase residential energy efficiency. Not least, appliance-level consumption information can give rise to new business models (e.g., providing cross-selling offers for non-energy-efficient devices). With a recognition rate of about 90% the results of our evaluation study confirm the suitability of the general scheme and encourage us to intensify further research.

Future work consists of deploying the system in various households to gather real-world data that allows for more in-depth evaluation of AppliSense. Based on these experiences, we plan to analyze the algorithm's dependency on the number of manually recorded signatures and to implement relevant refinements. This also includes accuracy improvements through the extension from one to three phases (which helps in case two appliances are turned on/off at the same time) and a module for auto-identification of hard-wired heating and cooling devices. In order to deal with edges detected in the load curve that do not yet correspond to an existing signature in the database, we focus on the application of clustering concepts that automatically classify these events (and once a certain probability is reached, verify the match by pushing a notification to the user interface). We also envision the possibility to upload appliance signatures to a community platform [18]. In the long term, we would like to investigate the possibility of building a larger appliance signature base. In addition, we are considering methods to derive occupancy state from electricity and appliance use data, in order to use this information in a smart heating control strategy [20].

ACKNOWLEDGEMENTS

REFERENCES

[1] Mattern, F., Staake, T., and Weiss, M. ICT for Green – How Computers Can Help Us to Conserve Energy. In Proc. e-Energy2010, ACM, 2010.

[2] Weiss, M., Graml, T., Staake, T., Mattern, F., and Fleisch, E. Handy feedback: Connecting smart meters with mobile phones. In Proc. MUM09, ACM, 2009.

[3] Hart, G.W. Nonintrusive appliance load monitoring. Proc. of the IEEE, 80(12):1870-1891, 1992.

[4] Norford, L.K. and Leeb, S.B. Non-intrusive electrical load monitoring in commercial buildings based on steady-state and transient load-detection algorithms. Energy and Buildings, 24(1): 51-64, 1996.

[5] Laughman, C. Power signature analysis. IEEE Power and Energy Magazine, 1(2):56-63, 2003.

[6] Cole, A. and Albicki, A. Data extraction for effective non-intrusive identification of residential power loads. In Proc. Instrumentation and Measurement Technology Conference, IMTC/98, 1998.

[7] Powers, B.S.J. and Margossian, B. Using a rule-based algorithm to disaggregate end-use load profiles from premise-level data. Trans. IEEE Computer Applications in Power, 4(2):42-47, 1991.

[8] Prudenzi, A. A neuron nets based procedure for identifying domestic appliances pattern-of-use from energy recordings at meter panel. In Proc. IEEE Power Engineering Society Winter Meeting, 2002.

[9] Ruzzelli, A.G., Nicolas, C., Schoofs, A., and O'Hare, G.M.P. Real-Time Recognition and Profiling of Appliances through a Single Electricity Sensor. In Proc. SECON 2010, 2010.

[10] Marceau, M.L. and Zmeureanu, R. Nonintrusive load disaggregation computer program to estimate the energy consumption of major end uses in residential buildings. Energy Conversion and Management, 41(1):1389-1403, 2000.

[11] Farinaccio, R.Z.L. Using a pattern recognition approach to disaggregate the total electricity consumption in a house into the major end-uses. Energy and Buildings, 30(3):245-259, 1999.

[12] Lam, H., Fung, G., and Lee, W. A Novel Method to Construct Taxonomy Electrical Appliances Based on Load Signatures. IEEE Trans. Consumer Electronics, 53(2):653-660, 2007.

[13] Srinivasan, D., Ng, W.S., and Liew, A.C. Neural-Network-Based Signature Recognition for Harmonic Source Identification. IEEE Transactions on Power Delivery, 21(1):398-405, 2006.

[14] Kolter, Z., Batra, S., and Ng, A.Y. Energy Disaggregation via Discriminative Sparse Coding. In Proc. Neural Information Processing Systems 2010.

[15] Patel, S., Robertson, T., Kientz, J., Reynolds, M., and Abowd, G. At the Flick of a Switch: Detecting and Classifying Unique Electrical Events on the Residential Power Line. In Proc. Ubicomp07, 2007.

[16] Gupta, S., Reynolds, M.S., and Patel, S.N. ElectriSense: Single-Point Sensing Using EMI for Electrical Event Detection and Classification in the Home. In Proc. UbiComp2010, 2010.

[17] Guinard, D., Trifa, V., Mattern, F., and Wilde, E. From the Internet of Things to the Web of Things: Resource Oriented Architecture and Best Practices. In: Dieter Uckelmann, Mark Harrison, Florian Michahelles: Architecting the Internet of Things. Springer, 2011.

[18] Weiss, M., Staake, T., Mattern, F., and Fleisch, E. PowerPedia – Changing Energy Usage with the Help of a Community-based Smartphone Application. Journal of Pervasive and Ubiquitous Computing, 2012.

[19] Schucany, W.R. Kernel smoothers: an overview of curve estimators for the first graduate course in nonparametric statistics. Statistical Science, 19(4):663-675, 2004.

[20] Kleiminger, W., Santini, S., and Weiss, M. Opportunistic Sensing for Smart Heating Control in Private Households. In Proc. CONET 2011.